

## CMPT 412 Project 3 Report

Ritika Goyal (301401516)

### Project 3: Object Detection, Semantic Segmentation, and Instance Segmentation


I am using two late days for assignment3

(Kaggle submission under name RitikaGoyal)

#### Part1: Object Detection

In this section, first I loaded the data from train.json and split it into training and validation set with 168 items in training and 30 items in validation. For the test set, the annotations were kept empty. I cached the data which saved time for loading it as it was just loaded once. Then I visualised the three random images from training\_set which are shown below:





## List of configs and modifications used:

These are the configs that I used for training:

```
[13] '''
# Set the configs for the detection part in here.
# TODO: approx 15 lines
'''

cfg = get_cfg()
cfg.OUTPUT_DIR = "{}output/".format(BASE_DIR)

cfg.merge_from_file(model_zoo.get_config_file("COCO-Detection/faster_rcnn_X_101_32x8d_FPN_3x.yaml"))
cfg.DATASETS.TRAIN = ("data_detection_train")
cfg.DATASETS.TEST = ()
cfg.DATALOADER.NUM_WORKERS = 2
# cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml")
cfg.SOLVER.IMS_PER_BATCH = 2
cfg.SOLVER.BASE_LR = 0.00025 # pick a good LR
cfg.SOLVER.MAX_ITER = 1000 #500
cfg.SOLVER.STEPS = [] # do not decay learning rate
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 512 # faster, and good enough for this toy dataset (default: 512)
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 1 # only has one class (ballon). (see https://detectron2.readthedocs.io/tutorials/datasets.html#to-use-the-coco-format)
# NOTE: this config means the number of classes, but a few popular unofficial tutorials incorrect uses num_classes + 1
# NOTE: this config means the number of classes, but a few popular unofficial tutorials incorrect uses num_classes + 1

os.makedirs(cfg.OUTPUT_DIR, exist_ok=True)
```


## Explanation:

Initially, I tested different learning rates to train the model. I tried 0.001, 0.0001 and 0.0005 which had similar issues that they started great for first few iterations and then converged really fast, and loss fluctuated. Choosing 0.00025 as learning rate improved the accuracy as this learning rate is smaller than others that I tried and large learning rate does not converge to local minima.

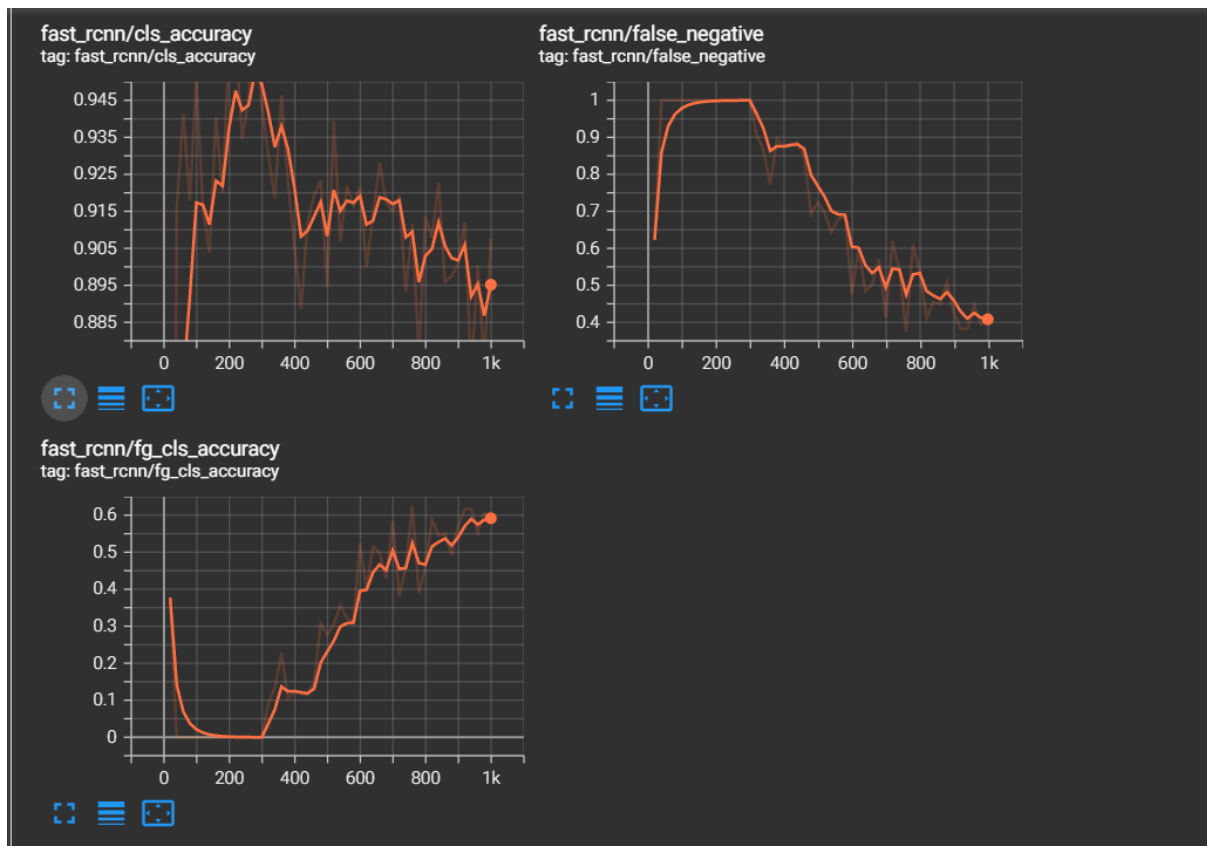
I used 1000 Max\_ITER. The initial given max\_iter was 500 but loss was still decreasing, So different values of max\_iter were tried until loss was stable.

The model that I used was faster\_rcnn\_x\_101\_32x8d\_FPN\_3x.yaml because it had higher train mem that is 6.7 GB and it converged better as compared to faster\_rcnn\_R\_101\_FPN\_3x.yaml and its box AP was also slightly higher. 2

## Plot for total training loss:




## Plot for accuracy:






## Visualised samples from test set and predicted results:

The visualised images using the above-mentioned configs and shown below:










### Visualisation of test samples when max\_iter = 500.







As it is seen, that there are so many planes that are not identified in test set when `max_iter` was 500, therefore `max_iter` = 1000 improved the performance significantly.

## **Part2: Semantic Segmentation**

### **Hyperparameters used were:**

- `batch_size` = 4
- `learning_rate` = 0.001
- `num_epochs` = 50
- `optimizer` = SGD optimizer

### **Network Architecture:**

I added four more down and up sampling layers in the network. The input layer was given (3,16) and then each down sampling layer incrd the output to the twice of input. After down sampling, the up sampling was done which decreases the output to half of the input. Making model like this, increased the feature channels which helped in detecting features. The optimizer was not changed and SGD optimizer was used. Number of epoch were made to 50 as loss became

stable. The loss at first epoch was 0.62 which got decreased to 0.16. The screenshots of the loss are given below:

```
+ Code + Text
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
Epoch: 45, Loss: 0.16584248840808868
100% ██████████ 1674/1674 [01:19<00:00, 22.60it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
Epoch: 46, Loss: 0.16459138691425323
100% ██████████ 1674/1674 [01:19<00:00, 22.83it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
Epoch: 47, Loss: 0.16249027848243713
100% ██████████ 1674/1674 [01:19<00:00, 21.50it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
Epoch: 48, Loss: 0.1615828424692154
100% ██████████ 1674/1674 [01:19<00:00, 21.39it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
Epoch: 49, Loss: 0.1602380871772766
```

The loss function is the default loss function, nn.BCEWithLogitsLoss()

**Final IoU:**

Mean IoU : 0.81

```
/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py:481:
cpuset_checked))
100% ██████████ 837/837 [00:23<00:00, 37.32it/s]
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:29: UserWarning: To copy construct from a tensor, it is recommended to

#images: 6694, Mean IoU: 0.812196756916907
```

**Visualised images from test set:**

**Image1: (original , ground truth, predicted)**



**Image2: (original , ground truth, predicted)**



**Image3: (original , ground truth, predicted)**



**Part3: Instance Segmentation:**

**Kaggle submission:**



Kaggle submission under the name RitikaGoyal

**Best Score on Kaggle: 0.60771**



Visualisation of results:

Image1: (original, ground truth, prediction)



**Image2: (original, ground truth, prediction)**







Image3: (original, ground truth, prediction)





## Part4: Mask R-CNN

### Visualization and the Evaluation:

Image1:



Image2:



Image3:



### **Observation:**

Part1 model works better than this one as it is seen in the visualisation as well that many planes are not detected by this model.

AP50 of part1 is 38.554 and part4 is 15.033. The mean IoU of part 3 is 0.81 but due to less AP50, it shows that this model is less accurate than previous one.

The screenshots are mentioned below:

## Part1:

```
[11/06 23:16:15 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.01 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.142
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.386
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.061
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.158
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.161
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.110
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.011
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.083
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.187
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.167
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.190
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.219
[11/06 23:16:15 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 14.213 | 38.554 | 6.053 | 15.787 | 16.074 | 11.011 |
OrderedDict([('bbox', {'AP': 14.213204817920975, 'AP50': 38.55438353822251, 'AP75': 6.052737722128967, 'APs':
```

## Part4:

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.018
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.044
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.252
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.004
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.028
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.053
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.024
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.065
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.289
[11/07 05:57:09 d2.evaluation.coco_evaluation]: Evaluation results for segm:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 3.979 | 15.033 | 1.173 | 1.774 | 4.441 | 25.165 |
OrderedDict([('bbox', {'AP': 16.838123079673075, 'AP50': 28.01757456538375, 'AP75': 17.776661
```



